

# Table of Contents

---

## Resource Management Lab

1. Manage Users, Projects, and Quotas
    - 1.1. Create Project
    - 1.2. View Resources in Web Console
    - 1.3. Apply Quota to Project
    - 1.4. Apply Limit Ranges to Project
    - 1.5. Test Quotas
    - 1.6. Run Pod
  2. Create Services and Routes
  3. Explore Containers
    - 3.1. Explore Route Container
    - 3.2. Explore Registry Container
  4. Create Persistent Volume for Registry
    - 4.1. Create NFS Export for Registry
    - 4.2. Attach Persistent Volume to Registry
- 

## Resource Management Lab

---

In this lab, you learn how to manage OpenShift Enterprise resources.

- **Manage Users, Projects, and Quotas**

In this section, you create projects and test the use of quotas and limits.

- **Create Services and Routes**

In this section, you manually create services and routes for pods and review the changes to a service when scaling an application.

- **Explore Containers**

In this section, you run commands within active pods and explore the

`docker-registry` and `Default Router` containers.

- **Create Persistent Volume for Registry**

In this section, you create a persistent volume for your registry, attach it to

`deploymentConfiguration`, and redeploy the registry.

---

## 1. Manage Users, Projects, and Quotas

---

### 1.1. Create Project

---

1. Connect to the **master00** host:

```
[root@oselab-GUID ~]# ssh master00-$guid
```

2. On the master host, run **oadm** to create and assign the administrative user **andrew** to a project:

```
[root@master00-GUID ~]# oadm new-project resourcemanagement --display-name="Resources Management" \
--description="This is the project we use to learn about resource management" \
--admin=andrew --node-selector='region=primary'
```



**andrew** can create his own project with the **oc new-project** command, an option you will experiment with later in this course. Note that defining the **--node-selector** is optional, especially since you already defined the default **node-selector** in a previous lab.

## 1.2. View Resources in Web Console

Now have a look at the web console, which has been completely redesigned for version 3.

1. Open your web browser and go to

**<https://master00-GUID.oslab.opentlc.com:8443>**.



The web console could take up to 90 seconds to become available after a restart of the master.

The first time you access the URL, you might need to accept the self-signed SSL certificate.

2. When prompted, type the username and password, as follows:

- **Username:** **andrew**
- **Password:** **r3dh4t1!**

3. In the web console, click the **Resources Management** project.



The project is empty because it has no data. You enter data in this part of the lab.

## 1.3. Apply Quota to Project

1. Create a quota definition file:

```
[root@master00-GUID ~]# cat << EOF > quota.json
{
```

```

"apiVersion": "v1",
"kind": "ResourceQuota",
"metadata": {
  "name": "test-quota"
},
"spec": {
  "hard": {
    "memory": "512Mi",
    "cpu": "20",
    "pods": "3",
    "services": "5",
    "replicationcontrollers": "5",
    "resourcequotas": "1"
  }
}
}
EOF

```

2. On the master host, do the following:

a. Run **oc create** to apply the file you just created:

```
[root@master00-GUID ~]# oc create -f quota.json --namespace=resourcemanagement
```

b. Verify that the quota exists:

```
[root@master00-GUID ~]# oc get -n resourcemanagement quota
```

NAME	AGE
test-quota	8s

c. Verify the limits and examine the usage:

```
[root@master00-GUID ~]# oc describe quota test-quota -n resourcemanagement
```

```

Name:      test-quota
Namespace: resourcemanagement
Resource   Used Hard
-----
cpu        0 20
memory     0 512Mi
pods       0 3
replicationcontrollers 0 5
resourcequotas 1 1
services   0 5

```

3. On the web console, click the **Resource Management** project.
4. Click the **Settings** tab for information on the quota.

## 1.4. Apply Limit Ranges to Project

For quotas to be effective, you must create *limit ranges*. They allocate the maximum, minimum, and default memory and CPU at both the pod and container level. Absent defaults for containers, projects with quotas fail because the deployer and other infrastructure pods are unbounded and, therefore, forbidden.

1. Create the **limits.json** file:

```
[root@master00-GUID ~]# cat << EOF > limits.json
{
  "kind": "LimitRange",
  "apiVersion": "v1",
  "metadata": {
    "name": "limits",
    "creationTimestamp": null
  },
  "spec": {
    "limits": [
      {
        "type": "Pod",
        "max": {
          "cpu": "500m",
          "memory": "750Mi"
        },
        "min": {
          "cpu": "10m",
          "memory": "5Mi"
        }
      },
      {
        "type": "Container",
        "max": {
          "cpu": "500m",
          "memory": "750Mi"
        },
        "min": {
          "cpu": "10m",
          "memory": "5Mi"
        },
        "default": {
          "cpu": "100m",
          "memory": "100Mi"
        }
      }
    ]
  }
}
```

EOF

2. On the master host, run `oc create` against the `limits.json` file and the `resourcemanagement` project:

```
[root@master00-GUID ~]# oc create -f limits.json --namespace=resourcemanagement
```

3. Review your limit ranges:

```
[root@master00-GUID ~]# oc describe limitranges limits -n resourcemanagement
```

```
Name: limits
Namespace: resourcemanagement
Type Resource Min Max Request Limit Limit/Request
-----
Pod memory 5Mi 750Mi - - -
Pod cpu 10m 500m - - -
Container memory 5Mi 750Mi 100Mi 100Mi -
Container cpu 10m 500m 100m 100m -
```

## 1.5. Test Quotas



You are running commands as the Linux users `andrew` and `root` in a lab environment. In a real-world scenario, users, would, of course, issue `oc` commands from their workstations and not from the OpenShift Master.

1. Authenticate to OpenShift Enterprise and choose your project:
  - a. Connect to the OpenShift Enterprise master according to the procedure you followed previously.
  - b. When prompted, type the username and password:
    - **Username:** `andrew`
    - **Password:** `r3dh4t1!`

```
[root@master00-GUID ~]# su - andrew
[andrew@master00-GUID ~]$ oc login -u andrew --insecure-skip-tls-verify --
server=https://master00-{guid}.oslab.opentlc.com:8443
```

- The output is as follows:

```
Login successful.
```

Using project "resourcemanagement".  
Welcome! See 'oc help' to get started.



This lab shows you the manual, step-by-step method of creating each object. There are easier ways to create a deployment and its components. One of those ways is the **oc new-app** command, which is covered later in this lab.

## 2. Create the **hello-pod.json** pod definition file:

```
[andrew@master00-GUID ~]$ cat <<EOF > hello-pod.json
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "hello-openshift",
    "creationTimestamp": null,
    "labels": {
      "name": "hello-openshift"
    }
  },
  "spec": {
    "containers": [
      {
        "name": "hello-openshift",
        "image": "openshift/hello-openshift:v1.0.6",
        "ports": [
          {
            "containerPort": 8080,
            "protocol": "TCP"
          }
        ],
        "resources": {
        },
        "terminationMessagePath": "/dev/termination-log",
        "imagePullPolicy": "IfNotPresent",
        "capabilities": {},
        "securityContext": {
          "capabilities": {},
          "privileged": false
        }
      }
    ],
    "restartPolicy": "Always",
    "dnsPolicy": "ClusterFirst",
    "serviceAccount": ""
  },
  "status": {}
}
```

EOF

## 1.6. Run Pod

Here, you create a simple pod without a *route* or *service*:

1. Create and verify the **hello-openshift** pod:

```
[andrew@master00-GUID ~]$ oc create -f hello-pod.json
pods/hello-openshift

[andrew@master00-GUID ~]$ oc get pods
NAME                READY    STATUS    RESTARTS   AGE
hello-openshift     1/1      Running   0           8s
```

2. Run **oc describe** for details on your pod:

```
[andrew@master00-GUID ~]$ oc describe pod hello-openshift
Name:      hello-openshift
Namespace:  resourcemanagement
Image(s):  openshift/hello-openshift:v1.0.6
Node:      node00-GUID.oslab.opentlc.com/192.168.0.200
Start Time:   Thu, 26 Nov 2015 21:23:27 -0500
Labels:      name=hello-openshift
Status:      Running
Reason:
Message:
IP:          10.1.2.2
Replication Controllers: <none>
Containers:
  hello-openshift:
    Container ID:
docker://e36321aabe1cb64e3da054128818dedd8ec3891dbf8aa758c72a96fc1180eee
    Image:      openshift/hello-openshift:v1.0.6
    Image ID:
docker://bba2117915baabfd05932dc916306bae2c51d15848592c3018e7af0308dee519
    QoS Tier:
      cpu: Guaranteed
      memory: Guaranteed
    Limits:
      cpu: 100m
      memory: 100Mi
    Requests:
      cpu: 100m
      memory: 100Mi
    State: Running
      Started:   Thu, 26 Nov 2015 21:23:32 -0500
    Ready:      True
    Restart Count: 0
    Environment Variables:
Conditions:
```

```

Type    Status
Ready   True
Volumes:
  default-token-rnadp:
    Type: Secret (a secret that should populate this volume)
    SecretName: default-token-rnadp
Events:
  FirstSeen    LastSeen    Count    From              SubobjectPath    Reason    Message
  -----
  4m          4m          1    {kubelet node00-GUID.oslab.opentlc.com}    implicitly required container
  POD Pulled   Container image "openshift3/ose-pod:v3.1.0.4" already present on
  machine
  4m          4m          1    {scheduler }      Scheduled Successfully assigned hello-openshift to
  node00-GUID.oslab.opentlc.com
  4m          4m          1    {kubelet node00-GUID.oslab.opentlc.com}    implicitly required container
  POD Created  Created with docker id f19fdc8fb3c8
  4m          4m          1    {kubelet node00-GUID.oslab.opentlc.com}    implicitly required container
  POD Started  Started with docker id f19fdc8fb3c8
  4m          4m          1    {kubelet node00-GUID.oslab.opentlc.com}    spec.containers{hello-
  openshift} Pulled Container image "openshift/hello-openshift:v1.0.6" already
  present on machine
  4m          4m          1    {kubelet node00-GUID.oslab.opentlc.com}    spec.containers{hello-
  openshift} Created Created with docker id e36321aabe1
  4m          4m          1    {kubelet node00-GUID.oslab.opentlc.com}    spec.containers{hello-
  openshift} Started Started with docker id e36321aabe1

```

### 3. Test that your pod is responding with **Hello OpenShift**:

```

[andrew@master00-GUID ~]$ ip=`oc describe pod hello-openshift|grep IP:|awk '{print $2}'`
[andrew@master00-GUID ~]$ curl http://${ip}:8080

```

- This output denotes a correct response:

```
Hello OpenShift!
```

### 4. Delete all the objects in your **hello-pod.json** definition file, which, at this point, is the pod only:

```
[andrew@master00-GUID ~]$ oc delete -f hello-pod.json
```



You can also delete a pod using the following command format:

```
oc delete pod hello-podname.
```

### 5. Create a new definition file that launches four **hello-openshift** pods:



```
[andrew@master00-GUID ~]$ cat << EOF > hello-many-pods.json
```

```
{
  "metadata":{
    "name":"quota-pod-deployment-test"
  },
  "kind":"List",
  "apiVersion":"v1",
  "items":[
    {
      "kind": "Pod",
      "apiVersion": "v1",
      "metadata": {
        "name": "hello-openshift-1",
        "creationTimestamp": null,
        "labels": {
          "name": "hello-openshift"
        }
      },
      "spec": {
        "containers": [
          {
            "name": "hello-openshift",
            "image": "openshift/hello-openshift:v1.0.6",
            "ports": [
              {
                "containerPort": 8080,
                "protocol": "TCP"
              }
            ],
            "resources": {
              "limits": {
                "cpu": "10m",
                "memory": "16Mi"
              }
            },
            "terminationMessagePath": "/dev/termination-log",
            "imagePullPolicy": "IfNotPresent",
            "capabilities": {},
            "securityContext": {
              "capabilities": {},
              "privileged": false
            }
          }
        ],
        "restartPolicy": "Always",
        "dnsPolicy": "ClusterFirst",
        "serviceAccount": ""
      },
      "status": {}
    },
    {
      "kind": "Pod",
      "apiVersion": "v1",
```

```

    "metadata": {
      "name": "hello-openshift-2",
      "creationTimestamp": null,
      "labels": {
        "name": "hello-openshift"
      }
    },
    "spec": {
      "containers": [
        {
          "name": "hello-openshift",
          "image": "openshift/hello-openshift:v1.0.6",
          "ports": [
            {
              "containerPort": 8080,
              "protocol": "TCP"
            }
          ],
          "resources": {
            "limits": {
              "cpu": "10m",
              "memory": "16Mi"
            }
          },
          "terminationMessagePath": "/dev/termination-log",
          "imagePullPolicy": "IfNotPresent",
          "capabilities": {},
          "securityContext": {
            "capabilities": {},
            "privileged": false
          }
        }
      ],
      "restartPolicy": "Always",
      "dnsPolicy": "ClusterFirst",
      "serviceAccount": ""
    },
    "status": {}
  },
  {
    "kind": "Pod",
    "apiVersion": "v1",
    "metadata": {
      "name": "hello-openshift-3",
      "creationTimestamp": null,
      "labels": {
        "name": "hello-openshift"
      }
    },
    "spec": {
      "containers": [
        {
          "name": "hello-openshift",
          "image": "openshift/hello-openshift:v1.0.6",

```

```
        "ports": [
          {
            "containerPort": 8080,
            "protocol": "TCP"
          }
        ],
        "resources": {
          "limits": {
            "cpu": "10m",
            "memory": "16Mi"
          }
        },
        "terminationMessagePath": "/dev/termination-log",
        "imagePullPolicy": "IfNotPresent",
        "capabilities": {},
        "securityContext": {
          "capabilities": {},
          "privileged": false
        }
      }
    ],
    "restartPolicy": "Always",
    "dnsPolicy": "ClusterFirst",
    "serviceAccount": ""
  },
  "status": {}
},
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "hello-openshift-4",
    "creationTimestamp": null,
    "labels": {
      "name": "hello-openshift"
    }
  },
  "spec": {
    "containers": [
      {
        "name": "hello-openshift",
        "image": "openshift/hello-openshift:v1.0.6",
        "ports": [
          {
            "containerPort": 8080,
            "protocol": "TCP"
          }
        ],
        "resources": {
          "limits": {
            "cpu": "10m",
            "memory": "16Mi"
          }
        }
      }
    ]
  }
}
```

```

    },
    "terminationMessagePath": "/dev/termination-log",
    "imagePullPolicy": "IfNotPresent",
    "capabilities": {},
    "securityContext": {
      "capabilities": {},
      "privileged": false
    }
  }
],
"restartPolicy": "Always",
"dnsPolicy": "ClusterFirst",
"serviceAccount": ""
},
"status": {}
}
]
}
EOF

```

6. Create the items in the **hello-many-pods.json** file:

```

[andrew@master00-GUID ~]$ oc create -f hello-many-pods.json
pod "hello-openshift-1" created
pod "hello-openshift-2" created
pod "hello-openshift-3" created
Error from server: Pod "hello-openshift-4" is forbidden: limited to 3 pods

```



Because you defined a quota before, **oc create** created three pods only instead of four.

7. Delete the object in the **hello-many-pods.json** definition file (the four pods):

```

[andrew@master00-GUID ~]$ oc delete -f hello-many-pods.json

```

8. (Optional) Create a project, set the quota with a pod value of **10**, and run **hello-many-pods.json**.

## 2. Create Services and Routes

1. As **andrew**, create a project called **scvslab**:

```

[andrew@master00-GUID ~]$ oc new-project scvslab --display-name="Services Lab" \
--description="This is the project we use to learn about services"

```

- The output looks like this:

```
Now using project "svclab" on server "https://master00-
GUID.oslab.opentlc.com:8443".
```



To switch between projects, run `oc project _projectname_`.

2. Create the `hello-service.json` file:

```
[andrew@master00-GUID ~]$ cat <<EOF > hello-service.json
{
  "kind": "Service",
  "apiVersion": "v1",
  "metadata": {
    "name": "hello-service",
    "labels": {
      "name": "hello-openshift"
    }
  },
  "spec": {
    "selector": {
      "name": "hello-openshift"
    },
    "ports": [
      {
        "protocol": "TCP",
        "port": 8888,
        "targetPort": 8080
      }
    ]
  }
}
EOF
```

3. Create the `hello-service` service:

```
[andrew@master00-GUID ~]$ oc create -f hello-service.json
service "hello-service" created
```

4. Display the services that are running in the current project:

```
[andrew@master00-GUID ~]$ oc get services
NAME                CLUSTER_IP          EXTERNAL_IP    PORT(S)    SELECTOR
AGE
hello-service       172.30.xxx.yyy      <none>         8888/TCP    name=hello-openshift
20s
```

5. Examine the details of your service. Note the following:
- **Selector:** Describes which pods the service selects or lists.
  - **Endpoints:** Displays all the pods that are currently listed (none in your current project).

```
[andrew@master00-GUID ~]$ oc describe service hello-service
Name:      hello-service
Namespace:  svcslab
Labels:     name=hello-openshift
Selector:   name=hello-openshift
Type:       ClusterIP
IP:         172.30.231.196
Port:       <unnamed> 8888/TCP
Endpoints:  <none>
Session Affinity: None
No events.
```

6. Create pods according to the **hello-many-pods.json** definition file:

```
[andrew@master00-GUID ~]$ oc create -f hello-many-pods.json
```

7. Wait a few seconds and check the service again.

- The pods that share the label **name=hello-openshift** are all listed:

```
[andrew@master00-GUID ~]$ oc describe service hello-service
Name:      hello-service
Namespace:  svcslab
Labels:     name=hello-openshift
Selector:   name=hello-openshift
Type:       ClusterIP
IP:         172.30.231.196
Port:       <unnamed> 8888/TCP
Endpoints:  <none>
Session Affinity: None
No events.

[andrew@master00-GUID ~]$ oc create -f hello-many-pods.json
pod "hello-openshift-1" created
pod "hello-openshift-2" created
pod "hello-openshift-3" created
pod "hello-openshift-4" created
[andrew@master00-GUID ~]$ oc describe service hello-service
Name:      hello-service
Namespace:  svcslab
Labels:     name=hello-openshift
Selector:   name=hello-openshift
Type:       ClusterIP
```

```
IP: 172.30.231.196
Port: <unnamed> 8888/TCP
Endpoints: 10.1.1.2:8080,10.1.1.3:8080,10.1.2.5:8080 + 1 more...
Session Affinity: None
No events.
```

## 8. Test that your service is working:

```
[andrew@master00-GUID ~]$ ip=`oc describe service hello-service|grep IP:|awk
'{print $2}'`
[andrew@master00-GUID ~]$ curl http://{ip}:8888
Hello OpenShift!
```

## 9. Expose your service with the **oc expose** command to create routes for your application:

```
[andrew@master00-GUID ~]$ oc expose service/hello-service --hostname=hello2-
openshift.cloudapps-${guid}.oslab.opentlc.com
```

## 10. View the route:

```
[andrew@master00-6b80 ~]$ oc get routes
```

NAME	HOST/PORT	PATH
SERVICE	LABELS	
hello-service	hello2-openshift.cloudapps-GUID.oslab.opentlc.com	
hello-service		

## 11. Test the route:

```
[andrew@master00-GUID ~]$ curl http://hello2-openshift.cloudapps-
${guid}.oslab.opentlc.com
Hello OpenShift!
```

---

# 3. Explore Containers

Next, take a look at the route and registry containers.

## 3.1. Explore Route Container

### 3.1.1. Create Applications As Examples

#### 1. As **andrew**, create a project called **explore-example**:

```
[andrew@master00-GUID ~]$ oc new-project explore-example --display-name="Explore
Example" \
```

```
--description="This is the project we use to learn about connecting to pods"
```

2. Applying the same image as before, run **oc new-app** to deploy **hello-openshift**:

```
[andrew@master00-GUID ~]$ oc new-app --docker-image=openshift/hello-openshift:v1.0.6 -l "todelete=yes"
--> Found Docker image 7ce9d7b (10 weeks old) from Docker Hub for "openshift/hello-openshift:v1.0.6"
* An image stream will be created as "hello-openshift:v1.0.6" that will track this image
* This image will be deployed in deployment config "hello-openshift"
* Ports 8080/tcp, 8888/tcp will be load balanced by service "hello-openshift"
--> Creating resources with label todelete=yes ...
ImageStream "hello-openshift" created
DeploymentConfig "hello-openshift" created
Service "hello-openshift" created
--> Success
Run 'oc status' to view your app.
```

3. Verify that **oc new-app** has created a pod and the service.

```
[andrew@master00-GUID ~]$ oc get service
NAME                                CLUSTER_IP          EXTERNAL_IP          PORT(S)              SELECTOR
AGE
hello-openshift    172.30.60.163      <none>               8080/TCP,8888/TCP
deploymentconfig=hello-openshift,todelete=yes    2m
[andrew@master00-GUID ~]$ oc get pods
NAME                                READY    STATUS    RESTARTS   AGE
hello-openshift-1-g3xow            1/1      Running   0           2m
```

4. Expose the service and create a route for the application:

```
[andrew@master00-GUID ~]$ oc expose service hello-openshift --
hostname=explore.cloudapps-${guid}.oslab.opentlc.com
```

5. In a later section, you explore the **docker-registry** container. To save time, start an S2I build now to push an image into the registry:

```
[andrew@master00-GUID ~]$ oc new-app https://github.com/openshift/sinatra-example -l "todelete=yes"
```

### 3.1.2. Connect to Default Router Container

1. As **root**, execute the **bash** shell inside the router with the **oc exec** command along with the default router's pod name. You have two options.

#### Option 1



```
[root@master00-GUID ~]# oc get pods
NAME                                READY    REASON    RESTARTS    AGE
docker-registry-2-snarn            1/1      Running    0            17h
trainingrouter-1-jm5zk            1/1      Running    0            18h
[root@master00-GUID ~]# oc exec -ti trainingrouter-1-jm5zk /bin/bash
```

## Option 2

```
[root@master00-GUID ~]# oc exec -ti `oc get pods | awk '/route/ { print $1; }`
"/bin/bash"
```

- With either option, this prompt is displayed:

```
[root@infranode00-GUID conf]#
```



You are now running **bash** inside the container. Also, the prompt specifies that you are on the **infranode** host. That is because the router container resolves the host name through the host's IP address.

## 2. Do the following:

- Run **id**.
- Run **pwd** and **ls** and note the directory you are in.
- Run **grep SERVERID** on the **haproxy.config** file.
- Run **cat haproxy.config** to verify that your configuration file is empty and then view the process status.

```
[root@infranode00-GUID conf]# id
uid=0(root) gid=0(root) groups=0(root)

[root@infranode00-GUID conf]# pwd
/var/lib/haproxy/conf

[root@infranode00-GUID conf]# ls
default_pub_keys.pem  os_edge_http_be.map    os_reencrypt.map
error-page-503.html  os_edge_http_expose.map  os_sni_passthrough.map
haproxy-config.template  os_edge_http_redirect.map  os_tcp_be.map
haproxy.config        os_http_be.map

[root@infranode00-GUID conf]# grep SERVERID haproxy.config
    cookie OPENSIFT_explore-example_hello-openshift_SERVERID insert indirect
nocache httponly
    cookie OPENSIFT_svcslab_hello-service_SERVERID insert indirect nocache
httponly
```

```
[root@infranode00-GUID conf]# ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
root          1        0  0 02:07 ?           00:00:14 /usr/bin/openshift-router
root        243        0  0 22:08 ?           00:00:00 /bin/bash
root        319        1  0 22:11 ?           00:00:00 /usr/sbin/haproxy -f /var/lib/
root        342       243  0 22:16 ?           00:00:00 ps -ef

[root@infranode00-GUID conf]# cat haproxy.config
```

e. Examine the output, which looks like this:

```
backend be_http_explore-example_hello-openshift

    mode http
    option redispatch
    option forwardfor
    balance leastconn
    timeout check 5000ms
    http-request set-header X-Forwarded-Host %[req.hdr(host)]
    http-request set-header X-Forwarded-Port %[dst_port]
    http-request set-header X-Forwarded-Proto https if { ssl_fc }

    cookie OPENSIFT_explore-example_hello-openshift_SERVERID insert indirect
    nocache httponly
    http-request set-header X-Forwarded-Proto http

    http-request set-header Forwarded for=%[src],host=%[req.hdr(host)],proto=%
[req.hdr(X-Forwarded-Proto)]

    server 10.1.1.7:8080 10.1.1.7:8080 check inter 5000ms cookie 10.1.1.7:8080

...
...
```

- Note the following:
  - The route is the one you created in the previous lab.
  - The route points to the endpoints directly.

3. As **andrew**, scale **hello-openshift** to have five replicas of its pod:

```
[andrew@master00-GUID ~]$ oc get deploymentconfig # or oc get dc
NAME          TRIGGERS          LATEST
hello-openshift  ConfigChange, ImageChange  1

[andrew@master00-GUID ~]$ oc scale dc hello-openshift --replicas=5
deploymentconfig "hello-openshift" scaled
```

4. Go back to the router container and view the **haproxy.config** file again:

```
[root@infranode00-GUID conf]# grep -A 25 backend.*explore-example_hello-openshift
haproxy.config
```

```
backend be_http_explore-example_hello-openshift
```

```
mode http
option redispatch
option forwardfor
balance leastconn
timeout check 5000ms
http-request set-header X-Forwarded-Host %[req.hdr(host)]
http-request set-header X-Forwarded-Port %[dst_port]
http-request set-header X-Forwarded-Proto https if { ssl_fc }
```

```
    cookie OPENSIFT_explore-example_hello-openshift_SERVERID insert indirect
nocache httponly
```

```
    http-request set-header X-Forwarded-Proto http
```

```
    http-request set-header Forwarded for=%[src],host=%[req.hdr(host)],proto=%
[req.hdr(X-Forwarded-Proto)]
```

```
server 10.1.1.7:8080 10.1.1.7:8080 check inter 5000ms cookie 10.1.1.7:8080
```

```
server 10.1.1.8:8080 10.1.1.8:8080 check inter 5000ms cookie 10.1.1.8:8080
```

```
server 10.1.1.9:8080 10.1.1.9:8080 check inter 5000ms cookie 10.1.1.9:8080
```

```
server 10.1.2.10:8080 10.1.2.10:8080 check inter 5000ms cookie 10.1.2.10:8080
```

```
server 10.1.2.11:8080 10.1.2.11:8080 check inter 5000ms cookie 10.1.2.11:8080
```

- All of your pods within the **haproxy** configuration are listed.



Remember, the router routes proxy connections to the pods directly and not through the service. The router uses the service only to obtain a list of the pod endpoints (IP addresses).

## 3.2. Explore Registry Container

Ensure that your build from earlier is complete.

1. As user **andrew**, run the following to see the build:

```
[andrew@master00-GUID ~]$ oc logs builds/sinatra-example-1
...
...
...
I1120 02:16:05.875303      1 sti.go:298] Successfully built
172.30.41.32:5000/svcslab/sinatra-example:latest
I1120 02:16:06.512944      1 cleanup.go:23] Removing temporary directory /tmp/s2i-
```

```
build079968192
I1120 02:16:06.513477      1 fs.go:99] Removing directory '/tmp/s2i-
build079968192'
I1120 02:16:06.546932      1 sti.go:213] Using provided push secret for pushing
172.30.41.32:5000/svcslab/sinatra-example:latest image
I1120 02:16:06.547064      1 sti.go:217] Pushing
172.30.41.32:5000/svcslab/sinatra-example:latest image ...
I1120 02:19:58.237018      1 sti.go:233] Successfully pushed
172.30.41.32:5000/svcslab/sinatra-example:latest
```



This step takes a while on the lab environment's hardware. If the build is not yet complete, feel free to take a quick break here.

2. As **root**, execute the **bash** shell inside the registry container by running **oc exec** along with the **docker-registry** pod name:

```
[root@master00-GUID ~]# oc exec -ti `oc get pods | awk '/registry/ { print $1; }'` /bin/bash
```

3. Do the following:

- Run **id**.
- Run **pwd** and **ls** and note the directory you are in.
- Run **cat config.yml** to verify that your configuration file is empty.

```
bash-4.2$ id
uid=1000000000 gid=0(root) groups=0(root)
bash-4.2$ pwd
/
bash-4.2$ ls
bin  config.yml  etc  lib  media  opt  registry  run  srv  tmp  var
boot  dev  home  lib64  mnt  proc  root  sbin  sys  usr
bash-4.2$ cat config.yml
version: 0.1
log:
  level: debug
http:
  addr: :5000
storage:
  cache:
    layerinfo: inmemory
  filesystem:
    rootdirectory: /registry
auth:
  openshift:
    realm: openshift
middleware:
  repository:
```

```
- name: openshift
bash-4.2$
```

#### 4. View the repositories and images that are available:

```
bash-4.2$ cd /registry/docker/registry/v2/repositories
bash-4.2$ ls
explore-example
bash-4.2$ ls explore-example/sinatra-example/_layers/
sha256
bash-4.2$ ls explore-example/sinatra-example/_layers/sha256/
50c4d0284685934ca2920fd6e056318cac1187773e8a239dd02d8f248a59d382
50de3644a809b46b344074ca0a691524eb06af3af6a07d25e90c25b50a00980f
9320560b540438b82b1bb1a51d035490812ad9298b945c041da3d0a4b646abf6
e1e04a46f510bf9b3fb68e6cf3fc027100cec875a7ff02e6d0da5206fa7f6b8c
```



Alternatively, if you configured persistent storage for your registry before, view the same in

```
/var/export/registry-storage/docker/registry/v2/.
```

#### 5. As user **andrew**, look at one of the pods you started earlier:

```
[andrew@master00-GUID ~]$ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-openshift-1-1ecah	1/1	Running	0	27m
hello-openshift-1-b8o3d	1/1	Running	0	27m
hello-openshift-1-g3xow	1/1	Running	0	45m
hello-openshift-1-rbfri	1/1	Running	0	27m
hello-openshift-1-yxidw	1/1	Running	0	27m
sinatra-example-1-build	0/1	Completed	0	11m
sinatra-example-1-yxyod	1/1	Running	0	8m

#### 6. Connect to the container:

```
[andrew@master00-GUID ~]$ oc exec -ti sinatra-example-1-yxyod "/bin/bash"
bash-4.2$
```

#### 7. Explore the container:

- Run **id**.
- Run **pwd** and **ls** and note the directory you are in.
- Run **ps -ef** to see what processes are running.

```
bash-4.2$ id
uid=1000050000 gid=0(root) groups=0(root)
```

```

bash-4.2$ pwd
/opt/app-root/src

bash-4.2$ ls
Gemfile      README.md  config.ru  example-mustache  public
Gemfile.lock app.rb     example-model  example-views  tmp
README      bundle    example-modular  example-views-modular

bash-4.2$ ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
1000050+      1         0  0  22:41 ?           00:00:01 ruby /opt/app-root/src/bundle/
1000050+     33         0  0  22:51 ?           00:00:00 /bin/bash
1000050+     62        33  0  22:51 ?           00:00:00 ps -ef

```



Your pod names and output differ slightly.

## 4. Create Persistent Volume for Registry

You learn in this lab how to create an NFS export for the registry and to attach the persistent volume to the registry.

### 4.1. Create NFS Export for Registry

1. As **root** on the **oselab** host, create a directory for your NFS export:

```

[root@oselab-GUID ~]# export volname=registry-storage
[root@oselab-GUID ~]# mkdir -p /var/export/pvs/${volname}
[root@oselab-GUID ~]# chown nfsnobody:nfsnobody /var/export/pvs/${volname}
[root@oselab-GUID ~]# chmod 700 /var/export/pvs/${volname}

```

2. Add this line to **/etc/exports**:

```

[root@oselab-GUID ~]# echo "/var/export/pvs/${volname} *(rw, sync, all_squash)" >>
/etc/exports

```

3. Restart NFS services:

```

[root@oselab-GUID ~]# systemctl restart rpcbind nfs-server nfs-lock nfs-idmap

```

4. As **root** on the **master** host, create a persistent volume-definition file named **registry-volume.json**:

```
[root@oselab-GUID ~]# ssh master00-$guid
[root@master00-GUID ~]# cat << EOF > registry-volume.json
{
  "apiVersion": "v1",
  "kind": "PersistentVolume",
  "metadata": {
    "name": "registry-storage"
  },
  "spec": {
    "capacity": {
      "storage": "15Gi"
    },
    "accessModes": [ "ReadWriteMany" ],
    "nfs": {
      "path": "/var/export/pvs/registry-storage",
      "server": "oselab-${GUID}.oslab.opentlc.com"
    }
  }
}

EOF
```

- In the **default** project, create the **registry-storage** persistent volume from the definition file:



You are creating the persistent volume in the **default** project because that is the project in which the registry runs.

```
[root@master00-GUID ~]# oc create -f registry-volume.json -n default
persistentvolume "registry-storage" created
```

- View the persistent volume you just created:

```
[root@master00-GUID ~]# oc get pv
```

NAME	LABELS	CAPACITY	ACCESSMODES	STATUS	CLAIM	REASON
AGE						
pv21	<none>	5Gi	RWO	Available		
20h						
pv22	<none>	5Gi	RWO	Available		
20h						
pv23	<none>	5Gi	RWO	Available		
20h						
registry-storage	<none>	15Gi	RWX	Available		
43s						

- Create a **registry-volume-claim.json** claim-definition file to claim your volume:

```
[root@master00-GUID ~]# cat << EOF > registry-volume-claim.json
{
  "apiVersion": "v1",
  "kind": "PersistentVolumeClaim",
  "metadata": {
    "name": "registry-claim"
  },
  "spec": {
    "accessModes": [ "ReadWriteMany" ],
    "resources": {
      "requests": {
        "storage": "15Gi"
      }
    }
  }
}

EOF
```

8. Create the **registry-claim** claim from the definition file:

```
[root@master00-GUID ~]# oc create -f registry-volume-claim.json -n default
persistentvolumeclaim "registry-claim" created
```

9. View the persistent volume you created, whose status is **Bound**:

```
[root@master00-GUID ~]# oc get pv
```

NAME	REASON	AGE	LABELS	CAPACITY	ACCESSMODES	STATUS	CLAIM
pv21		20h	<none>	5Gi	RWO	Available	
pv22		20h	<none>	5Gi	RWO	Available	
pv23		20h	<none>	5Gi	RWO	Available	
registry-storage-claim		2m	<none>	15Gi	RWX	Bound	default/registry-claim

10. View the persistent volume claim you created, whose status is also **Bound**:

```
[root@master00-GUID ~]# oc get pvc
```

NAME	AGE	LABELS	STATUS	VOLUME	CAPACITY	ACCESSMODES
registry-claim	43s	<none>	Bound	registry-storage	15Gi	RWX

## 4.2. Attach Persistent Volume to Registry



1. Assuming that your registry is already running, obtain the names of

**deploymentConfigurations**:

```
[root@master00-GUID ~]# oc get dc
NAME                TRIGGERS    LATEST
docker-registry     ConfigChange 1
trainingrouter      ConfigChange 1
```

2. Run **oc volume** to modify **DeploymentConfiguration**.
3. Add the **registry-storage** volume to the registry's **DeploymentConfiguration**, hence redeploying the registry:

```
[root@master00-GUID ~]# oc volume dc/docker-registry --add --overwrite -t
persistentVolumeClaim \
--claim-name=registry-claim --name=registry-storage
```

4. Run **oc get pods**:

```
[root@master00-GUID ~]# oc get pods
NAME                READY    STATUS    RESTARTS   AGE
docker-registry-2-d9niy 1/1      Running   0           31s
trainingrouter-1-xcz9o  1/1      Running   0           21h
```



Along with the deletion of the first **docker-registry** container, all the images it stored were also deleted. Now that your registry contains a persistent volume, images are saved even if you delete or replace the **docker-registry** pod.

5. As **andrew** on the **master** host, start an application based on the **<https://github.com/openshift/sti-php>** repository that would require an S2I build:

```
[root@master00-GUID ~]# su - andrew
[andrew@master00-GUID ~]$ oc new-app
openshift/php-https://github.com/openshift/sti-php -l "todelete=yes"
--> Found image 355eabc (2 weeks old) in image stream "php in project openshift"
under tag :latest for "openshift/php"
    * A source build using source code from https://github.com/openshift/sti-php
will be created
    * The resulting image will be pushed to image stream "sti-php:latest"
    * This image will be deployed in deployment config "sti-php"
    * Port 8080/tcp will be load balanced by service "sti-php"
--> Creating resources with label todelete=yes ...
ImageStream "sti-php" created
BuildConfig "sti-php" created
```

```
DeploymentConfig "sti-php" created
Service "sti-php" created
--> Success
Build scheduled for "sti-php" - use the logs command to track its progress.
Run 'oc status' to view your app.
```

6. Check the build logs to ensure that the build is complete and has been pushed into the registry:

```
[andrew@master00-GUID ~]$ oc logs -f builds/sti-php-1
I1126 23:24:28.604316      1 sti.go:298] Successfully built
172.30.42.118:5000/default/sti-php:latest
I1126 23:24:28.716843      1 cleanup.go:23] Removing temporary directory /tmp/s2i-
build491090638
I1126 23:24:28.717016      1 fs.go:99] Removing directory '/tmp/s2i-
build491090638'
I1126 23:24:28.740315      1 sti.go:213] Using provided push secret for pushing
172.30.42.118:5000/default/sti-php:latest image
I1126 23:24:28.740431      1 sti.go:217] Pushing 172.30.42.118:5000/default/sti-
php:latest image ...
I1126 23:25:51.808905      1 sti.go:233] Successfully pushed
172.30.42.118:5000/default/sti-php:latest
```



The **-f** flag sets **oc logs** to "follow" the log, similar to **tail -f**.

7. On the NFS server, **oselab**, verify that the registry is using the **registry-storage** volume:

```
[root@oselab-GUID ~]# find /var/export/pvs/registry-storage | grep sti-php
... Omitted output ...
... Omitted output ...
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-php/_uploads
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-php/_layers
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-php/_layers/sha256
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/812413b2241fa8ff63cb2747bf62e516ff4dc953b1332014faa551655c0ed608
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/812413b2241fa8ff63cb2747bf62e516ff4dc953b1332014faa551655c0ed608
/link
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/b18d4a50300b72f417496313920eff6d4bad00c0f1446686e3d5f157d255d0d2
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
```

```
example/sti-
php/_layers/sha256/b18d4a50300b72f417496313920eff6d4bad00c0f1446686e3d5f157d255d0d2
/link
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/50c4d0284685934ca2920fd6e056318cac1187773e8a239dd02d8f248a59d382
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/50c4d0284685934ca2920fd6e056318cac1187773e8a239dd02d8f248a59d382
/link
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/9320560b540438b82b1bb1a51d035490812ad9298b945c041da3d0a4b646abf6
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_layers/sha256/9320560b540438b82b1bb1a51d035490812ad9298b945c041da3d0a4b646abf6
/link
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-php/_manifests
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-php/_manifests/revisions
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-php/_manifests/revisions/sha256
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_manifests/revisions/sha256/5b8677660e3f1959a0eb44f1ac87200329c721ff4acd8c59f78
a8d0afa5dd425
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_manifests/revisions/sha256/5b8677660e3f1959a0eb44f1ac87200329c721ff4acd8c59f78
a8d0afa5dd425/signatures
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_manifests/revisions/sha256/5b8677660e3f1959a0eb44f1ac87200329c721ff4acd8c59f78
a8d0afa5dd425/signatures/sha256
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_manifests/revisions/sha256/5b8677660e3f1959a0eb44f1ac87200329c721ff4acd8c59f78
a8d0afa5dd425/signatures/sha256/561fd3acac303de8a9c4de202a2e3169bb47f5c03586358d13d
374832e983df5
/var/export/pvs/registry-storage/docker/registry/v2/repositories/explore-
example/sti-
php/_manifests/revisions/sha256/5b8677660e3f1959a0eb44f1ac87200329c721ff4acd8c59f78
a8d0afa5dd425/signatures/sha256/561fd3acac303de8a9c4de202a2e3169bb47f5c03586358d13d
374832e983df5/link
... Omitted output ...
/var/export/pvs/registry-storage/docker/registry/v2/blobs/sha256/53
/var/export/pvs/registry-
storage/docker/registry/v2/blobs/sha256/53/53aca6d1d55ccf8f9074725396099dc9592641a2
ae233cb8b1b2de2c800410cb
/var/export/pvs/registry-
storage/docker/registry/v2/blobs/sha256/53/53aca6d1d55ccf8f9074725396099dc9592641a2
ae233cb8b1b2de2c800410cb/data
/var/export/pvs/registry-storage/docker/registry/v2/blobs/sha256/b1
```

```
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/b1/b18d4a50300b72f417496313920eff6d4bad00c0  
f1446686e3d5f157d255d0d2  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/b1/b18d4a50300b72f417496313920eff6d4bad00c0  
f1446686e3d5f157d255d0d2/data  
/var/export/pvs/registry-storage/docker/registry/v2/blobs/sha256/50  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/50/50c4d0284685934ca2920fd6e056318cac118777  
3e8a239dd02d8f248a59d382  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/50/50c4d0284685934ca2920fd6e056318cac118777  
3e8a239dd02d8f248a59d382/data  
/var/export/pvs/registry-storage/docker/registry/v2/blobs/sha256/93  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/93/9320560b540438b82b1bb1a51d035490812ad929  
8b945c041da3d0a4b646abf6  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/93/9320560b540438b82b1bb1a51d035490812ad929  
8b945c041da3d0a4b646abf6/data  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/93/931b7ebd6c92756356ae4174a02b845480c5c548  
84875533ffa4cbef3872199a  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/93/931b7ebd6c92756356ae4174a02b845480c5c548  
84875533ffa4cbef3872199a/data  
/var/export/pvs/registry-storage/docker/registry/v2/blobs/sha256/81  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/81/812413b2241fa8ff63cb2747bf62e516ff4dc953  
b1332014faa551655c0ed608  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/81/812413b2241fa8ff63cb2747bf62e516ff4dc953  
b1332014faa551655c0ed608/data  
/var/export/pvs/registry-storage/docker/registry/v2/blobs/sha256/56  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/56/561fd3acac303de8a9c4de202a2e3169bb47f5c0  
3586358d13d374832e983df5  
/var/export/pvs/registry-  
storage/docker/registry/v2/blobs/sha256/56/561fd3acac303de8a9c4de202a2e3169bb47f5c0  
3586358d13d374832e983df5/data
```



You can see that previously created images are not in the registry, they were created before the registry was restarted and given a persistent volume.